

Zebu: A Tool for Specifying Reversible LALR(1) Parsers

Joachim Laubsch (laubsch@hplabs.hp.com)

September 16, 2007

Application Engineering Department
Software Technology Laboratory
Hewlett-Packard Laboratories
1501 Page Mill Road, Bldg. 1U-17
P.O. Box 10490
Palo Alto, Calif. 94304-1126 laubsch@hpl.hp.com

(415) 857-7695

Contents

1	Introduction	2
2	The Representations of Grammars in Files	4
2.1	Grammar notation	4
2.1.1	Grammar Rules	4
3	Grammar Options	8
3.1	Keyword Arguments to Grammar Construction	8
3.2	Defining a Domain	9
4	The Zebu Meta Grammar	12
4.1	Domain Definition	14
4.2	Example Grammars	14
4.3	The Kleene * Notation	17
5	Using the Compiler	18
5.1	Compiling a grammar	18
5.2	Loading a grammar	21
5.3	Parsing a string with a grammar	21
5.4	Parsing from a file with a grammar	22
5.5	Parsing from a list of tokens	23
5.6	Debugging a grammar	24
6	Lexical Analysis	24
6.1	Customization and Regular Expressions	24
6.2	Introducing new Categories by Regular Expressions	24
6.3	The functional interface to the parsing engine	26
7	Future Work	27

Abstract

Zebu¹ is part of a set of tools for the translation of formal languages. Zebu contains a LALR(1) parser generator like Yacc does. Aside from generating a parser, Zebu will also generate the inverse of a parser (unparser). In contrast to Yacc, the semantics is not given in terms of “routines” but declaratively in terms of typed feature structures.

The ability to declaratively define a reversible grammar, together with a rewrite-rule mechanism (Zebu-RR) for transforming abstract syntax trees constitute the basic tools for specifying translators for formal languages.

Keywords Formal language, LALR-grammar, parsing, translation, generation, interoperability, LEX, YACC.

1 Introduction

Our goal is to develop an environment for the design, analysis and manipulation of formal languages, such as programming languages, markup languages, data interchange formats or knowledge representation languages (such as the translation to and from KIF) [4]. Being able to design, analyze, and manipulate formal languages is crucial for achieving software interoperability [3], automatic code analysis, indexing, and retrieval for potential reuse. Zebu has been applied to writing translators for formal languages [6]. The main idea of this work is that a module m communicates by sending or receiving messages in some language $L(m)$, and that for various reasons different modules use different languages. For communication to be successful, translators have to be used. Zebu provides tools to define translators at a high level of abstraction².

McCarthy introduced the notion of “abstract” and “concrete” syntax. The concrete syntax describes the surface form of a linguistic expression, while the abstract syntax describes a (composite) object. E.g. “1+a” is the surface string rendered by a particular concrete syntax for an object described by an abstract syntax: an addition

¹“**zebu***n.*, *pl.* **-bus**, **-bu**: 1. an oxlike domestic animal (*Bos indicus*) native to Asia and parts of Africa: it has a large hump over the shoulders, short, curving horns, pendulous ears, and a large dewlap and is resistant to heat and insect-born diseases.” [Webster’s New World Dictionary.] A zebu should not be confused with a yacc or a gnu although it bears similarity to each of them.

²The rewrite-rule mechanism (Zebu-RR) is implemented, and will be described in a future report.

operation with two operands, the first being the numeral “1”, and the second being the variable named “a”.

Manipulation of linguistic expressions is much easier to express in the abstract syntax than in the concrete syntax.

If we were to design an algorithm for simplifying expressions of some language — say “arithmetic” — we would use as the front end the “arithmetic-parser” to translate into abstract syntax, then express the simplification rules in terms of tree transformation rules that operate on the abstract syntax, and finally add as the back-end the “arithmetic-unparser”.

More generally, if we were to design an algorithm for translating from language A to language B, we would define reversible grammars for languages A and B, and sets of rewrite rules to transform the abstract syntax trees from the domain of language A to the domain of language B. The front end would be the “A-parser” and the back-end the “B-unparser”

The work described in this report owes a lot to the pioneering research at Kestrel [9] that resulted in the **Refine**³ program transformation system [8]. The basic ideas underlying **Zebu** are already present in **Refine**. **Zebu** is much more compact than **Refine**⁴, and the semantics is expressed in typed feature structures. **Zebu** also offers the possibility of defining a meta-grammar. **Zebu** lacks **Refine**’s ability to declaratively specify transformations using a pattern language.⁵

The LALR(1) parsing table generated by **Zebu** follows algorithms described in [1] or [2]. The current implementation was developed from the **Scheme** program developed by William Wells and is written in **Common Lisp**.

The next section will explain how a grammar can be defined, and how semantics can be associated with a grammar rule. Section 3 describes the definition of the semantic domain. With this capability it is possible to state declaratively what the abstract syntax should look like. Section 4 describes a simpler grammar notation that is very close to ordinary BNF. Section 5 summarizes the functional interface of **Zebu** and explains how a parser can be customized. Section 6 describes how lexical analysis can be extended using regular expressions and parameterization.

³**Refine** is a trademark of Reasoning Systems, Palo Alto.

⁴**Zebu** runs on a MacIntosh in MacIntosh Common Lisp.

⁵**Zebu** can be obtained via anonymous ftp from ftp.cs.cmu.edu as a compressed tar file: /user/ai/lang/lisp/code/zebu/zebu-???.tar.gz. It contains several example grammar definitions.

2 The Representations of Grammars in Files

2.1 Grammar notation

We first describe the null-grammar, which is a powerful but verbose way to specify a grammar. Only a parser and optionally a domain will be generated but an unparser (printer) will not. If this is desired, you must use the notation of the meta-grammar "zebu-mg" which is described in section 4.

Non-terminals are represented by symbols, terminals (also referred to as keywords) by strings. There are the following open classes of non-terminals⁶:

`identifier` ::= $\langle \text{lisp symbol} \rangle$

`number` ::= $\langle \text{integer} \rangle$

`keyword` ::= $\langle \text{string} \rangle$

where

$\langle \text{integer} \rangle$::= $\langle \text{digit} \rangle^*$

$\langle \text{string} \rangle$::= " $\langle \text{character} \rangle^*$ "

A $\langle \text{lisp symbol} \rangle$ may be qualified by a package name, e.g. `zb:cons-1-3` is a valid identifier. In case packages should be disallowed during lexical analysis, the variable `*disallow-packages*` should be bound to *true*. (It defaults to *false*). The alphabetic case of a keyword is not significant if the variable `*case-sensitive*` is *false* (the default) when the grammar is loaded.

If alphabetic case of identifiers is to be preserved, `*preserve-case*` should be set to *true*. Other categories can be defined as regular expressions (see 6.2).

2.1.1 Grammar Rules

Grammar Rule Syntax A grammar file consists of a header (the "options list", see section 3.1) followed by one or more domain definitions or grammar rules. The non-terminal defined by the first grammar rule is also the *start-symbol* of the grammar. A parser will accept exactly the strings that rewrite to the *start-symbol*.

This example shows how a BNF-like rule can be encoded as a Zebu grammar rule (using the null-grammar):

- BNF rule example

$$\langle A \rangle ::= \langle B \rangle \mid \langle C \rangle \langle \text{number} \rangle \mid \text{"foo"} \langle A \rangle \mid \text{"c"} \mid \langle \text{the-empty-string} \rangle$$

⁶The Kleene * indicates 0 or more occurrences of the preceding constituent

- Zebu null-grammar example:

```
(defrule A
  := B                      ; (1)
  := (C NUMBER)            ; (2)
  := ("foo" A)             ; (3)
  := "c"                   ; (4)
  := ()                    ; (5)
)
```

The rule describes 5 productions, all deriving the non-terminal **A**. Each of the productions has the left-hand side **A**. The right-hand side of (1) consists of just one constituent, the non-terminal **B**. (2) has a right-hand of length 2, and its second constituent is the non-terminal **NUMBER** (which rewrites to any integer, real or rational). (3) is a recursive production. (4) contains just the terminal (or keyword) **"c"**. (5) derives the empty string.

None of these productions has a semantic action attached. By default, the semantic action is the **identity** function if the right-hand side of the rule consists of a single constituent and the **identity*** function otherwise. (**identity*** is defined as the function that returns all its arguments as a list.)

Grammar Rule Semantic Actions If we want to attach other than these default semantic actions, we have to use a **:build** clause after a production.

The build clause has the syntax:

```
<build clause> ::= :build (<lisp function> <argument list>)
<build clause> ::= :build <atomic lisp form>
<build clause> ::= :build (:form <lisp form>)
<build clause> ::= :build (:type <struct-type>
                           :map (((<non-terminal> . <Slot>)*)))
```

The first case

```
:build (<lisp function> <argument list>)
```

is like a function call. It may contain free variable occurrences. These will be bound to the non-terminal constituents of the same name occurring in the right-hand side of the production at the time of applying the semantic action.

In the second case

`:build` \langle atomic lisp form \rangle

the \langle atomic lisp form \rangle must be a function. It will be applied to the constituents of the right-hand side. This function should have the same number of arguments as the right-hand side of the corresponding production has constituents.

Since it happens often, that only some of the constituents of the right-hand side are selected, or combined, a few useful semantic actions have been predefined in **Zebu**.⁷

An example for such a predefined action is the function `cons-2-3` which takes 3 arguments and returns a *cons* of its second and third argument.

The third form of the `:build` clause is just a long way to write the first form, i.e.

`:build` (\langle lisp function \rangle \langle argument list \rangle)

is the same as

`:build` (`:form` (\langle lisp function \rangle \langle argument list \rangle))

Similarly,

`:build` (`progn` \langle atomic lisp form \rangle)

is the same as

`:build` (`:form` \langle atomic lisp form \rangle)

The last `:build` clause is more interesting:

`:build` (`:type` \langle struct-type \rangle
 `:map` (((\langle Nonterminal \rangle . \langle Slot \rangle)*))

where \langle struct-type \rangle is a symbol that must be the name of a structure type⁸. Instead of having to write the semantic action as a constructing form, we just have to specify the type and the mapping of non-terminals to slots, as in the following example⁹:

⁷These semantic actions (`cons-1-3` `cons-2-3` `empty-seq` `empty-set` `k-2-1` `k-2-2` `k-3-2` `k-4-3` `identity*` `seq-cons` `set-cons`) are described in the file "zebu-actions.lisp".

⁸a type defined by `defstruct` or `defclass`.

⁹(taken from the grammar named "pc1"; see the file "pc1.zb" in the test directory)

```

(defrule Boolean-Expr
  := (Formula.1 "and" Formula.2)
  :build (:type Boolean-And
          :map ((Formula.1 . :-rand1)
                (Formula.2 . :-rand2)))

  := (Formula.1 "or" Formula.2)
  :build (:type Boolean-Or
          :map ((Formula.1 . :-rand1)
                (Formula.2 . :-rand2)))
)

```

The map indicates that the slot `-rand1` is to be filled by the value of the non-terminal `Formula.1`, etc.

This example also makes use of the `".n"` notation: If on the right-hand side of a production a nonterminal occurs repeatedly, we distinguish it by appending `"."` and a digit, to the nonterminal (e.g. `Formula.1`).

The function `print-actions` applied to the name of a grammar may be used to find out what the generated code for the semantic actions looks like, e.g. after compiling the sample grammar `'pc1.zb'`:

```

(print-actions "pc1")

...
Rule: BOOLEAN-EXPR
(LAMBDA (FORMULA.1 DUMMY FORMULA.2)
  (DECLARE (IGNORE DUMMY))
  (MAKE-BOOLEAN-AND :-RAND1 FORMULA.1 :-RAND2 FORMULA.2))
(LAMBDA (FORMULA.1 DUMMY FORMULA.2)
  (DECLARE (IGNORE DUMMY))
  (MAKE-BOOLEAN-OR :-RAND1 FORMULA.1 :-RAND2 FORMULA.2))
...

```

These semantic actions have been generated from the `:build` clauses of the above rule for `Boolean-Expr`.

3 Grammar Options

3.1 Keyword Arguments to Grammar Construction

Some global information to control grammar compilation, lexical analysis, and the generation of semantic actions is declared in the beginning of a grammar file¹⁰. A grammar file must begin with a list of alternating keywords and arguments. The following keywords are valid:

¹⁰A grammar file has the default type ".zb".

<code>:name</code>	a string, the name of the grammar to be defined.
<code>:package</code>	a string, the name of the package where the non-terminal symbols and the function symbols used in semantic actions reside.
<code>:identifier-start-chars</code>	a string. During lexical analysis any character in this string can start an identifier non-terminal. The default is *identifier-start-chars* .
<code>:identifier-continue-chars</code>	a string. During lexical analysis any character in this string can continue an identifier (i.e. characters not in this string terminate identifier). The default is *identifier-continue-chars* .
<code>:intern-identifier</code>	<i>true</i> , if the identifier is to be returned as an interned Lisp symbol, or <i>false</i> if the identifier is to be returned as a string (default <i>true</i>).
<code>:string-delimiter</code>	a character, the character that delimits a string to be represented as a Common Lisp string. (default <code>#\"</code>)
<code>:symbol-delimiter</code>	a character, the character that delimits a string to be represented as a Common Lisp symbol. (default <code>#\'</code>)
<code>:domain</code>	a list, representing the type hierarchy of the domain. See section 3.2 below.
<code>:domain-file</code>	a string naming the file where the generated Common Lisp program that implements the domain will be stored. Definitions of functions for semantic actions and regular expression for lexical categories are kept here as well. This string defaults to the concatenation of the grammar's <code>:name</code> and <code>"-domain"</code> .
<code>:grammar</code>	a string, by default: "null-grammar" , naming the grammar to be used to parse the grammar defined in this file. If the grammar "zebu-mg" is used, an unparser will also be generated.
<code>:lex-cats</code>	an association list of terminal category names and regular expressions (see section 6.2).
<code>:white-space</code>	a list of characters each of which will be ignored before a token, (default <code>(#\Space #\Newline #\Tab)</code>)
<code>:case-sensitive</code>	<i>true</i> if the case of keywords is significant, <i>false</i> otherwise (default <i>false</i>).

3.2 Defining a Domain

The `:domain` keyword is used to specify a type hierarchy. This specification will expand into **defstruct** forms that implement this hierarchy. It is also possible to write such structure definitions directly into the grammar file. The argument to the

:domain keyword argument must be a list of the following form:

```

(⟨Root Struct⟩
 :subtype ⟨Struct Desc⟩
 :subtype ⟨Struct Desc⟩
 ...)

⟨Root Struct⟩ ::= ⟨Symbol⟩

⟨Struct Desc⟩ ::= ⟨Symbol⟩ |
                  ( ⟨Symbol⟩ :slots (⟨Slot⟩*) ) |
                  ( ⟨Symbol⟩ :slots (⟨Slot⟩*)
                    :subtype ⟨Struct Desc⟩
                    :subtype ⟨Struct Desc⟩
                    ... )

⟨Slot⟩ ::= ⟨Symbol⟩ | ( ⟨Slot Name⟩ ⟨Filler Type⟩ )
⟨Filler Type⟩ ::= ⟨Symbol naming type⟩

```

This describes the syntax for declaring a type hierarchy with root node $\langle \text{Root Struct} \rangle$. A node of the hierarchy tree can have children, denoted by **:subtype** followed by the structure description of the child node. Each node can have slots, described as a list following **:slots**. A child node inherits the slots of its parent node. The value of a slot can be type-restricted to $\langle \text{Filler Type} \rangle$.

$\langle \text{Root Struct} \rangle$ will be implemented as a structure type directly below the predefined structure type **kb-domain**, i.e. $(\text{kb-domain-p } x)$ is *true* for any instance of a subtype of $\langle \text{Root Struct} \rangle$. **kb-domain** is the top of the domain hierarchy.

The type **kb-sequence** is already predefined as a subtype of **kb-domain**. It has the slots **first** and **rest**.

Similarly, types **number**, **string**, and **identifier** are predefined as subtypes of **kb-domain**.

Two objects of type **kb-domain** can be compared for equality with the functions **kb-equal** and **kb-compare**.

kb-equal *a b*

function

a and *b* are assumed to be of type **kb-domain**. If they are **equal** they are also **kb-equal**. But in contrast to **equal** it is possible to define which slots are to be examined by **kb-equal** when comparing the components of *a* and *b*. These relevant slots are called *tree attributes*, and the macro **def-tree-attributes** is used to define

these for a particular type. The rationale for having this equality relation is that it is often useful to store comments or auxiliary information with the feature structures produced by parsing.

In feature structures the value of a relevant feature (or slot) may be declared to be a set (using `def-tree-attributes`). If a slot has been declared set-valued, the `kb-equal` comparison will use set equality for values of that slot (represented as lists).

`def-tree-attributes` *type slot1 slot2 ..* *macro*

`def-tree-attributes` defines *slot1 slot2 ...* as tree attributes for instances of type *type*.

If *slot* is a symbol, this symbol is defined as a tree attribute. Otherwise *slot* must be of the form (*symbol* :set). As before, the *symbol* becomes a tree-attribute, and furthermore it is declared set-valued.

Example domain definition The grammar defined in “pc1.zb” accepts a simple propositional calculus language with sentences such as

`walks(agent: John),`

which yields the following abstract syntax (printed out using the Common Lisp structure printer):

```
#S(ATOMIC-WFF -PREDICATE WALKS
  -ROLE-ARGUMENT-PAIRS #S(ROLE-ARGUMENT-PAIR
    -ROLE AGENT
    -ARGUMENT JOHN) )
```

The types — such as `ATOMIC-WFF` and `ROLE-ARGUMENT-PAIR` — are defined by the following domain declaration:

```
:domain (PC ;; PC is the root type of the hierarchy
  :subtype (Formula
    :subtype (Propositional-variable :slots (-name))
    :subtype (Boolean-Expr
      :slots ((-rand1 Formula)
        (-rand2 Formula))
      :subtype Boolean-Or
      :subtype Boolean-And))
  :subtype (Boolean-Op :slots (-name))
```

```

:subtype (Atomic-Wff
          :slots (-predicate
                  (-Role-Argument-Pairs KB-Sequence)))
:subtype (Role-Argument-Pair :slots (-Role -Argument))
)

```

Note the use of the predefined type `KB-Sequence`. It is used to construct the list of `Role-Argument-Pairs` in the following rule:

```

(defrule Role-Argument-Pairs
  := ()

  := (Role-Argument-Pair Role-Argument-Pairs)
  :build (:type KB-Sequence
          :map ((Role-Argument-Pair . :first)
                (Role-Argument-Pairs . :rest)))
)

```

4 The Zebu Meta Grammar

Using `"zebu-mg"` as the `:grammar` argument in the grammar options indicates that the following grammar is to be preprocessed with the grammar `"zebu-mg"` before compilation.

The advantages of the meta-grammar (versus the default null-grammar) are a more concise representation of rules, automatic generation of the functions that implement the semantic actions and reversibility of the grammar (generation of printing functions – the unparser).

The disadvantage of using `"zebu-mg"` is that the semantics is limited to constructing typed feature structures. But these have great expressive power, and furthermore could subsequently be transformed into some other program. Typed feature structures are ideally suited to present abstract syntax. The fact that unification, specialization and generalization are well defined operations on feature structures, makes them appropriate for further transformations (by e.g. `Zebu-RR`). For an introduction into feature structures see [5].

Since there is a restricted way of expressing the semantics of a rule – namely as a typed feature structure – the grammar compiler will be able to generate code for the domain hierarchy and print-functions associated with each type of that domain.

”zebu-mg” is defined in terms of the null-grammar described above¹¹.

BNF description of “zebu-mg”:

$\langle \text{Zebu-Grammar} \rangle$	$::= \langle \text{Options} \rangle \langle \text{Domain-Defn} \rangle^* \langle \text{zb-rule} \rangle$
$\langle \text{Domain-Defn} \rangle$	$::= \langle \text{Type-name} \rangle \text{ ":=" } \langle \text{Feat-Term} \rangle$ $\quad [\text{ "<<" "print-function:" Identifier ">>" }] \text{ ";" }$
$\langle \text{zb-rule} \rangle$	$::= \langle \text{Non-terminal} \rangle \text{ "-->" } \langle \text{Rhs} \rangle \text{ ";" }$
$\langle \text{Rhs} \rangle$	$::= \langle \text{Rhs1} \rangle \langle \text{More-Rhs} \rangle \mid \langle \text{Kleene-Rhs} \rangle$
$\langle \text{Rhs1} \rangle$	$::= \langle \text{Constituent} \rangle^* [\text{ "{" } \langle \text{Semantics} \rangle \text{ "}" }]$
$\langle \text{Constituent} \rangle$	$::= \langle \text{Identifier} \rangle \mid \langle \text{String} \rangle$
$\langle \text{More-Rhs} \rangle$	$::= \mid \langle \text{Rhs1} \rangle \langle \text{More-Rhs} \rangle$
$\langle \text{Semantics} \rangle$	$::= \langle \text{Feat-Term} \rangle$

A $\langle \text{Feat-Term} \rangle$ is a typed attribute value matrix.

$\langle \text{Feat-Term} \rangle$	$::= [\langle \text{Type-name} \rangle \text{ ":" }] \langle \text{Conj} \rangle$
$\langle \text{Conj} \rangle$	$::= [\text{ "[" } \langle \text{Label-value-pair} \rangle \text{ " * " }] \text{ " }$
$\langle \text{Label-value-pair} \rangle$	$::= [\text{ "(" } \langle \text{Identifier} \rangle \langle \text{Feat-Term} \rangle \text{ ")" }]$
$\langle \text{Type-name} \rangle$	$::= \langle \text{Identifier} \rangle$

$\langle \text{Options} \rangle$ is described in section 3.

This BNF-notation makes use of

1. star (*) for 0 or more repetitions of the preceding constituent
2. bar (|) for alternation
3. brackets ([]) for marking the enclosed constituents as optional
4. a quotation symbol (") for delimiting keywords

The above definition is somewhat oversimplified, since it does not deal with the “.n” notation for $\langle \text{Constituent} \rangle$: if on the right-hand side of a production a non-terminal occurs repeatedly, we can distinguish the occurrences by appending “.n” and a digit to the identifier. The semantics can then unambiguously refer to an occurrence of a constituent.

The semantics is described as a typed feature structure. Names of variables occurring in feature term position correspond to constituent names in the right-hand

¹¹You may study the definition of the meta grammar in terms of the null-grammar in the file “zebu-mg.zb”.

side of the rule. The effect of applying a rule is to instantiate a feature structure of the type described in the rule semantics, substituting variables with their values.

If the relation between semantics and syntax is one-to-one, the inverse of a parser, a printer, can be generated.

4.1 Domain Definition

Although it is possible to specify the hierarchy of domain types using the `:domain` keyword as in section 3.2, a more convenient syntax is offered by the meta above grammar rule `<Domain-Defn>`.

The type definition

$$atype := super: [(s_1) \dots (s_n)];$$

will define the type *atype* inheriting from *super*, and having slots s_1 through s_n .

$$atype := [(s_1) \dots (s_n)];$$

is as above but defines the type *atype* as a subtype of the top type named **kb-domain**.

A slot may be type restricted as in:

$$atype := super: [(s_1 \text{ KB-sequence})];$$

which restricts s_1 to be of type **KB-sequence**. An optional *print-function* may be specified, as in

$$atype := super: [(s_1)] << \text{print-function: } print-atype >>;$$

Here we supply for *atype* its own printer called *print-atype* and no printer will be generated for *atype*. Usually it is not necessary to provide a print-function, but if the grammar is ambiguous, this is a way to force a particular canonical unparser.

4.2 Example Grammars

Example Grammar for Arithmetic Expressions

```
(:name "arith-exp" :grammar "zebu-mg")
```

```
;; Domain definition
```

```

Arith-exp := Kb-domain: [];
Factor    := Arith-exp: [(-value)] <<print-function: Print-factor>>;
Mult-op   := Arith-exp: [(-arg1) (-arg2)];
Plus-op   := Arith-exp: [(-arg1) (-arg2)];

;; Productions

EE --> EE "+" TT { Plus-op: [(-arg1 EE) (-arg2 TT)] }
      | TT ;

TT --> TT "*" F   { Mult-op: [(-arg1 TT) (-arg2 F)] }
      | F ;

F -->  "(" EE ")"      { factor: [(-value EE)] }
      | IDENTIFIER     { factor: [(-value IDENTIFIER)] }
      | NUMBER         { factor: [(-value NUMBER)] } ;

```

The semantics of the first rule says that an object of type `+-op` should be created with slot `-arg1` filled with the value of `EE` and `-arg2` filled with the value of `TT`.

Example Grammar for Propositional Calculus This grammar defines the same domain as above (3.2). Compiling it generates a parser and a generator.

```

(:name "pc2"
 :package "CL-USER"
 :grammar "zebu-mg")

;; Domain definition

Formula := kb-domain: [];

Propositional-variable := Formula: [(-name) ];
P-Formula              := Formula: [(-content) ];
Boolean-Expr           := Formula: [(-rand1 Formula) (-rand2 Formula)];
  Boolean-Or            := Boolean-Expr: [];
  Boolean-And           := Boolean-Expr: [];
Atomic-Wff             := Formula: [(-predicate)
                                (-Role-Argument-Pairs kb-sequence)];

```



```

Role-Argument-Pair := kb-domain: [(-Role) (-Argument)];

;; Productions

Formula --> Propositional-variable
          | Boolean-Expr
          | "(" Formula ")" {P-Formula:[(-content Formula)]}
          | Atomic-Wff;

Propositional-Variable
--> Identifier {Propositional-variable: [(-name Identifier)]};

Boolean-Expr --> Formula.1 "and" Formula.2
                {Boolean-And: [(-rand1 Formula.1)
                              (-rand2 Formula.2)]}

                | Formula.1 "or" Formula.2
                {Boolean-Or: [(-rand1 Formula.1)
                              (-rand2 Formula.2)]};

Atomic-Wff --> Identifier "(" Role-Argument-Pairs ")"
              { Atomic-Wff:
                [(-predicate Identifier)
                 (-Role-Argument-Pairs Role-Argument-Pairs)]};

Role-Argument-Pairs -->
  | Role-Argument-Pair Role-Argument-Pairs
  { RAP-list: [(-first Role-Argument-Pair)
               (-rest Role-Argument-Pairs)]};

Role-Argument-Pair -->
  Identifier ":" Term
  {Role-Argument-Pair: [(-Role Identifier)
                       (-Argument Term)]};

Term --> Identifier | Number ;

```

4.3 The Kleene * Notation

The meta-grammar “zebu-mg” provides an abbreviated notation for repeated occurrences of a non-terminal, separated by a keyword. The syntax for this is:

$$\begin{aligned}\langle \text{Kleene-Rhs} \rangle & ::= \langle \text{Identifier} \rangle * \langle \text{String} \rangle & (1) \\ \langle \text{Kleene-Rhs} \rangle & ::= \langle \text{Identifier} \rangle + \langle \text{String} \rangle & (2)\end{aligned}$$

The meaning of (1) is that 0 or more occurrences of the constituent named by $\langle \text{Identifier} \rangle$ and separated by $\langle \text{String} \rangle$ will be accepted by this rule, and that the sequence of the results of these constituents will be returned as the semantics of $\langle \text{Kleene-Rhs} \rangle$. The meaning of (2) is the same, except that at least one occurrence of the constituent has to be found.

The semantics of a $\langle \text{Kleene-Rhs} \rangle$ production is an implicit kb-sequence construction. The Kleene-constituent ($\langle \text{Identifier} \rangle$ concatenated with $*$ or $+$) is bound in the semantics of the production, e.g.

```
Disjunction --> Conjunction+ "|"
              {Disj: [(-terms Conjunction+)]};
```

builds a structure of type `Disj` with the `-terms` slot filled by the value of the Kleene-constituent `Conjunction+`.

Example grammar using Kleene * Notation

```
(:name "mini-la" :grammar "zebu-mg" )

;; Domain definition

Program := [(-stmts kb-sequence)];
Application := [(-function) (-args kb-sequence)];

;; rules

Program --> "begin" Stmt+ ";" "end"
          { Program: [(-stmts Stmt+)] } ;

Stmt     --> Identifier | Appl | Program ;
```

```

Appl    --> Identifier "(" Arg* " " ")"
          {Application: [(-function Identifier) (-args Arg*)]};

Arg      --> Identifier | Number | Appl ;

```

Compiling this grammar generates a parser/unparser (i.e. the printing routines are generated automatically).

```

(zb:read-parser "begin A; B ; C end"
  :grammar (zb:find-grammar "mini-la"))

```

returns a structure of type PROGRAM which is printed in the syntax of “mini-la”:

```

begin A;B;C end
> (describe *)
begin A;B;C end is a structure of type PROGRAM.
It has 1 slot, with the following values:
-STMTS:                A;B;C

(describe (PROGRAM--STMTS *))
A;B;C is a structure of type KB-SEQUENCE.
It has 2 slots, with the following values:
FIRST:                  A
REST:                   B C

```

5 Using the Compiler

5.1 Compiling a grammar

The Zebu-compiler¹² can be called using any of the functions: `zebu-compile-file`, `compile-slr-grammar`, `compile-lalr1-grammar`.

```

zebu-compile-file                                     function
               grammar-file &key output-file grammar verbose

```

¹²For installation see appendix A.

This compiles the LALR(1) grammar in a file named *grammar-file*. The *output-file* defaults to a file with the same name as *grammar-file* but type "tab". The grammar used for compilation defaults to the null-grammar. If *verbose* is *true*, conflict warnings will be printed. `zebu-compile-file` returns the pathname of *output-file*.

Example:

```
(let ((*warn-conflicts* t)
      (*allow-conflicts* t))
  (zebu-compile-file "dangelse.zb"
                    :output-file "/tmp/dangelse.tab"))

; Zebu Compiling (Version 2.0)
; "~/zebu/test/dangelse.zb" to "/tmp/dangelse.tab"

Reading grammar from dangelse.zb

Start symbols is: S

4 productions, 8 symbols
.....9 item sets
.....
.....
;;; Warning: ACTION CONFLICT!!!-- state: 8
;;;          old entry: (6 :S 2)  new entry: (6 :R 2)
;;;
;;; Warning: Continuing to build tables despite conflicts...
;;;          Will prefer old entry: (6 :S 2)

Dumping parse tables to /tmp/dangelse.tab
#P"/tmp/dangelse.tab"
```

`*warn-conflicts*`

variable

If *true* during LALR-table construction, shift-reduce conflicts will be reported. By default, `*warn-conflicts*` is *false*.

`*allow-conflicts*`

variable

If *true* during LALR-table construction, shift-reduce conflicts will be resolved in favor of the old entry. By default, `*allow-conflicts*` is *true*.

`*check-actions*`

variable

If *true* the semantic action associated with a production will be compiled at grammar compilation time in order to display possible warning messages. By default, `*check-actions*` is *false*.

`compile-slr-grammar grammar-file &key output-file grammar`

function

This is like `zebu-compile-file`, but an SLR-table will be made.

Example:

```
(compile-slr-grammar "dangelse.zb"
  :output-file "/tmp/dangelse.tab")
```

Reading grammar from dangelse.zb

Start symbols is: S

4 productions, 8 symbols
.....9 item sets

Dumping parse tables to /tmp/dangelse.tab
#P"/tmp/dangelse.tab"

`compile-lalr1-grammar grammar-file &key output-file grammar`

function

This is like `zebu-compile-file`, but does not expand logical pathnames.

Example:

```
(compile-lalr1-grammar "dangelse.zb"
  :output-file "/tmp/dangelse.tab")
```

Reading grammar from dangelse.zb

Start symbols is: S

4 productions, 8 symbols
.....9 item sets

.....
.....

Dumping parse tables to /tmp/dangelse.tab
#P"/tmp/dangelse.tab"

5.2 Loading a grammar

`zebu-load-file` *filename* &key *verbose* *function*

filename should be the name of a compiled grammar file, i.e. a file of type "tab". If such a file can be found, it will be loaded, returning the grammar object needed for parsing. In case a domain-file was generated by compiling the grammar, it will also be loaded. The type of the domain-file is the first for which a file named *filename-domain.<type>* exists, by examining the lists

```
*load-binary-pathname-types* and  
*load-source-pathname-types*
```

for *.<type>* in turn.

The keyword argument *verbose* defaults to *true*.

Example:

```
(zebu-load-file "/tmp/dangelse.tab")  
<Zebu Grammar: dangelse>
```

It is possible to have many grammars loaded concurrently. Given the name of a grammar, one can find a grammar that has been loaded by:

`find-grammar` *name* *function*

name must be a string. If a grammar of the same name (ignoring case) has been loaded, the grammar object is returned, else *false* is returned.

Example:

```
(find-grammar "dangelse")  
<Zebu Grammar: dangelse>
```

5.3 Parsing a string with a grammar

`read-parser` *function*
string &key *grammar* *junk-allowed* *print-parse-errors* *error-fn* *start*

The argument of the `:grammar` keyword defaults to `*current-grammar*` (initially bound to the null-grammar), e.g.

```
(read-parser <string> :grammar (find-grammar <name>))
```

is equivalent to

```
(setq zebu:*current-grammar* (find-grammar <name>))
(read-parser <string>)
```

`read-parser` parses the string starting at the position indicated by `:start` (default 0).

`read-parser` takes the keyword argument `:junk-allowed`, which if *true* will return as second value an index to the unparsed remainder of the string in case not the entire string was consumed by the parse.

The keyword `:junk-allowed` has the same meaning as in the Common Lisp function `read-from-string`.

`:print-parse-errors` controls the printing of errors during parsing and defaults to *true*.

`:error-fn` is a function used to report errors, it defaults to the Common Lisp error function.

Example:

```
(read-parser "if f then if g then h else i"
             :grammar (find-grammar "dangelse"))
("if" F "then" ("if" G "then" H "else" I))

(read-parser "1 + a" :grammar (find-grammar "ex1"))
(+OP (EXPRESSION (TERM (FACTOR 1)))
      (TERM (FACTOR A)))
```

5.4 Parsing from a file with a grammar

`file-parser` *file* &key *grammar print-parse-errors verbose* *function*

`file-parser` parses expressions using the grammar specified by `:grammar`, reading from *file*. It returns a list of the parse-results, i.e. a list of what would have been returned by `read-parser`.

The `:grammar` argument defaults to `*current-grammar*` – which initially is bound to the “null-grammar”.

`:print-parse-errors` controls the printing of errors during parsing and defaults to *true*.

`:verbose` controls whether printing of parse-results occurs, and defaults to *true*.

The processing of comments by `file-parser` can be influenced by the following variables:

- `*comment-brackets*` is a list of bracket pairs. Everything between any of bracket pairs is ignored. Initially `*comment-brackets*` is set to:

```
(("#\" . "\"#)).
```

- `*comment-start*` A line beginning with this character is ignored. Initially `*comment-start*` is set to the semicolon character: `#\;`

Example:

```
(file-parser "sample-ex1" :grammar (find-grammar "ex1"))  
...
```

5.5 Parsing from a list of tokens

`list-parser` *token-list* &key *grammar* *junk-allowed* *function*

`list-parser` is like `read-parser` except that the tokens that are passed by the scanner to the driver are already given as the elements of *token-list*. This function is useful if the options for controlling lexical analysis given in section 3.1 are insufficient.

Example:

```
(let ((*current-grammar* (find-grammar "ex1")))  
  (list-parser '(1 "+" x "*" y)))  
(+OP (EXPRESSION (TERM (FACTOR 1)))  
  (*-OP (TERM (FACTOR X)) (FACTOR Y)))
```


5.6 Debugging a grammar

`debug-parser` &key *grammar lexer*

function

`debug-parser` will cause a trace of the parser to be displayed. The *grammar* keyword defaults to *true* and *lexer* defaults to *false*. If *lexer* is *true*, more information about lexical analysis (see section 6 below) will be displayed.

6 Lexical Analysis

6.1 Customization and Regular Expressions

It should only seldomly be necessary to write a lexical analyzer. Before you attempt to introduce your own lexical categories, check whether the following variables and keywords would suffice to parameterize lexical analysis:

```
*comment-start*
*comment-brackets*
*disallow-packages*
*preserve-case*
*case-sensitive*
:case-sensitive
:identifier-start-chars
:identifier-continue-chars
:string-delimiter
:symbol-delimiter
:white-space
:lex-cats
```

The lexical analyzer works in a top-down one token look-ahead way. It tries only to recognize tokens that would be legal continuations of the string parsed so far. In case lexical categories overlap this will serve to disambiguate tokenization.

6.2 Introducing new Categories by Regular Expressions

The keyword `:lex-cats` takes as argument an association list of the form:

```
(((<Category> <Regular Expression>)) *)
```

⟨Category⟩ is a symbol naming a lexical category and ⟨Regular Expression⟩ is a string representing a regular expression as defined in the GNU Emacs Lisp Manual [7]. The regular expression will be compiled into a Common Lisp function and invoked by `read-parser` before the built-in categories (Identifier, String, Number) are examined. The categories can be used in grammar rules like any of the built-in categories.

The regular expression compiler¹³ handles the following constructs:

. Period matches any single character except a newline.

* repeats preceding regular expression as many times as possible.

+ like * but must match at least once.

? like * but must match once or not at all.

[...] '[' begins a character set, which is terminated by ']'.

Character ranges can be indicated, e.g. a-z, 0-9.

[^...] forms the complement character set.

\$ matches only at the end of a line.

\(...\) is a grouping construct.

\⟨digit⟩ means: accept the same string as was matched by the group in position ⟨digit⟩.

Example:

```
:lex-cats ((BibTeX-in-braces "{[^\\n]}*"))
```

defines a new category `BibTeX-in-braces` which matches anything starting with "{", ending in "}", and not containing either a newline or "\\".

```
:lex-cats
((Ratio_Number "-?[0-9]+/[0-9]+")
 (Simple_Float "-?[0-9]*\\. [0-9]+"))
```

defines the syntax for rationals and floating point numbers. Note that the period needs to be escaped, since it is a special character of the regular expression language.

¹³Thanks to Lawrence E. Freil who wrote the main part of the Regular Expression Compiler.

6.3 The functional interface to the parsing engine

In case the above parameterization facilities for lexical analysis are insufficient or you want to use an existing lexical analyzer, you need to understand the functional interface to the parsing engine as implemented by the `lr-parse`.

`lr-parse` *function*
next-sym-fn error-fn grammar &optional *junk-allowed last-pos-fn*

`lr-parse` returns the result of parsing the token stream produced by *next-sym-fn* with *grammar* by the LALR(1) method. In case *junk-allowed* is *true* it produces as second value a handle to the yet unconsumed token stream by calling the function *last-pos-fn*.

next-sym-fn should be bound to a generator function — a function of no arguments — that will be called to produce the next token. It should return two values: (1) the token found and (2) the category of the token (obtained by the function `categorize`).

error-fn is the function to be called in case of an error. *grammar* is the grammar object that contains important information for lexical analysis, (e.g. the table of keywords).

To understand the interface to `lr-parse`, consider how `list-parser` (described above) might have been defined:

```
(defun list-parser (token-list &key (grammar *current-grammar*)
                    junk-allowed)
  (let ((last-position token-list)
        token1)
    (check-type token-list list)
    (lr-parse
     ;; The LEXER supplied to the parsing engine:
     #'(lambda ()
         (if (null token-list)
             (end-of-tokens-category grammar)
             (progn
              (setq last-position token-list
                    token1 (pop token-list))
              (categorize token1 grammar))))))
     ;; The error function supplied to the parsing engine:
     #'(lambda (string)
         (error "~S~% Remaining tokens: ~S~{ ~S~}"
                string token1 token-list)))
    grammar
    junk-allowed
    ;; Function that returns the remaining unparsed token-list
    #'(lambda () last-position))))
```

`end-of-tokens-category` *grammar*

function

`end-of-tokens-category` returns two values: a token signifying the end of the token stream and the appropriate lexical category.

`categorize` *token grammar*

function

`categorize` returns the *token* itself and its category, a number that represents one of `number`, `identifier`, `string` or a terminal token defined by `:lex-cats`.

7 Future Work

Translation involves three processes:

- parsing
- transformation
- generation

Zebu is a tool that helps in 1 and 3. There are cases where 2 reduces to the identity function, since the abstract syntax is the same for the source and the target language of translation. Examples for these “syntactic variants” are infix and prefix notation for arithmetic or boolean expressions.

In general, the situation is more complicated. For languages with the same expressive power, some transformation process can be defined. Between languages with different expressive power such a transformation is not always possible. For a language that is not Turing complete, it is not possible to express every computation, e.g. SQL cannot express recursion, and hence it is not possible to express the “ancestor” relation (which is recursively defined). A technique to represent transformation are “rewrite rule systems”. The **Refine** language [8] contains a rewrite-rule mechanism in which the rules are in terms of patterns of the concrete syntax. We have implemented a rewrite-rule system based on typed feature structures, called **Zebu-RR**, which will be described in a future report.

A Installation

There are two ways to install **Zebu**:

- Installation using **defsystem**

This makes it easier to load and compile grammars, since one does not need to remember the location of a module in a directory structure and the particular compilation and loading functions. To install, follow the directions in **ZEBU-sys.lisp**. You need the portable **defsys** for that. This is available as **Defsys.tar.gz** at the same place as **zebu-???.tar.gz**.

The file **ZEBU-sys.lisp** is used to load or compile **Zebu**, which actually consists of two systems (defined by **defsystem**)

Zebu	the runtime system
Zebu-compiler	the compiler

- Installation without **defsystem**

If you don’t want to use **defsystem**, load the file **COMPILE-ZEBU.lisp**, which

compiles the Zebu files in the right order. After loading the file `ZEBU-init.lisp` you can call:

`(zb:zebu)` to load the runtime system

or

`(zb:zebu-compiler)` to load the grammar compiler.

References

- [1] A.V. Aho and J.D. Ullman. *Principles of Compiler Design*. Addison Wesley, New York, 1979.
- [2] Charles N. Fischer and Richard J. LeBlanc. *Crafting a Compiler*. Benjamin/Cummings, Menlo Park, CA, 1988.
- [3] Michael R. Genesereth. An agent-based framework for software interoperability. Technical Report Logic-92-02, Department Of Computer Science, Stanford University, Stanford, 1992.
- [4] Michael R. Genesereth, Richard Fikes, et al. Knowledge interchange format, version 3.0. reference manual. Report Logic-92-1, Logic Group Report, Computer Science Department, Stanford University, Stanford, June 1992.
- [5] Mark Johnson. *Attribute Value Logic and the Theory of Grammar*. Center for the Study of Language and Information, Stanford, 1988.
- [6] Joachim Laubsch and Derek Proudian. A case study in REFINE: interfacing modules via languages. Report HPL-STL-TM-88-11, Hewlett Packard, 1988.
- [7] Bill Lewis, Dan LaLiberte, and the GNU Manual Group. *GNU Emacs Lisp Reference Manual*. The Free Software Foundation, Cambridge, MA, December 1990.
- [8] Reasoning Systems, Palo Alto, 3260 Hillview Ave., CA 94304. *Refine User's Guide*, 1989.
- [9] Douglas R. Smith, Gordon B. Kotik, and Stephen J. Westfold. Research on knowledge-based software environments at KESTREL institute. *IEEE Transactions on Software Engineering*, SE-11:1278–1295, November 1985.

Index

- *allow-conflicts*, 18
- *case-sensitive*, 5, 24
- *check-actions*, 18
- *comment-brackets*, 23, 24
- *comment-start*, 23, 24
- *disallow-packages*, 5, 24
- *preserve-case*, 5, 24
- *warn-conflicts*, 18
- :build semantic action, 6
- :case-sensitive, 9, 24
- :domain, 8
- :domain-file, 8
- :grammar, 8
- :identifier-continue-chars, 8, 24
- :identifier-start-chars, 8, 24
- :intern-identifier, 8
- :lex-cats, 8, 24
- :name, 8
- :package, 8
- :print-parse-errors, 22
- :string-delimiter, 8, 24
- :symbol-delimiter, 8, 24
- :white-space, 9, 24

- categorize, 27
- compile-lalr1-grammar, 20
- compile-slr-grammar, 20

- debug-parser, 23
- def-tree-attributes, 11
- domain
 - defining, 9, 13
 - top type, 10
- feature structures, 12
- file-parser, 22
- find-grammar, 21

- grammar
 - name, 8
 - options, 8

- kb-compare, 10
- kb-domain, 10
- kb-domain-p, 10
- kb-equal, 10
- kb-sequence, 10
- Kleene *, 5, 16, 17

- lexical category, 24
- list-parser, 23, 26
- lr-parse, 25

- meta grammar, 4, 12

- non-terminal, 5
 - “.n” notation, 8
- null-grammar, 4

- option list, 8

- print-actions, 8
- print-function, 12, 14

- read-parser, 17, 21
- regular expression, 9, 23

- semantic actions, 6, 7
- start-symbol, 5

- zebu, 28
- zebu-compile-file, 18
- zebu-compiler, 28
- zebu-load-file, 20
- zebu-mg, 12